



CRAQ Garbage Collection Bug


We describe a very minor bug in CRAQ's garbage collection. Consider the CRAQ deployment with three replicas shown below, where r_0 is the head and r_2 is the tail. Initially, all three replicas store a single version (i.e. version 0) of object x , which has value a.




A client sends a write request with value b to the head, r_0 . r_0 creates a new version of x and forwards the write request to r_1 , which also creates a new version of x .



Then, a client sends a read request of value x to the head. r_0 has more than one version of x , so it cannot perform the read locally. Instead, it sends a version request to the tail r_2 . When r_2 receives the version request, it queries to find that its version of x is 0. It sends 0 back to r_0 , but this message is delayed in network and is not yet received by r_0 .



Next, r_1 forwards the write request to the tail, and the tail successfully propagates write acknowledgements backwards through the chain. When r_0 and r_1 receive the acknowledgement, they garbage collect version 0 of x , as described in the CRAQ paper: **“When an acknowledgment message for an object version arrives at a node, the node marks the object version as clean. The node can then delete all prior versions of the object.”**



Finally, r_0 receives the delayed message, which includes version 0, from r_2 . r_0 then attempts to read version 0 of x . The CRAQ paper says **“by construction, the node is guaranteed to be storing this version of the object”**, but this is not true. r_0 does not have this version (actually, no node has this version), and is stuck.

